

**CSE 415 – Introduction to Parallel Computing**  
**Spring 2018, Homework 2**  
Due 5 pm, Friday, Feb 9<sup>th</sup>

*This homework is 15% of your **homework grade** (not your total grade).*

**Important note:** Please use a word processing software (e.g., MS Word, Mac Pages, Latex, etc.) to type your homework. Follow the submission instructions at the end to turn in an electronic copy of your work.

- 1) [25 pts] Performance Modeling: Assume you are evaluating a polynomial of the form,  
 $x = y^2 + z^3 + yz$ , which can be implemented as follows:

```
float x[N], y[N], z[N];  
  
for (i=0; i < N; ++i)  
    x[i] = y[i]*y[i] + z[i]*z[i]*z[i] + y[i]*z[i];
```

Here i is an integer and x,y and z are single precision floating point arrays.

- a. [10 pts] What is the arithmetic intensity of this kernel?
- b. [10 pts] Assume this kernel is to be executed on a processor that has 30 GB/sec of memory bandwidth. Under what conditions will this kernel be memory bound, and under what conditions will it be compute bound?
- c. [5 pts] Develop a roofline model for this processor, assuming it has a peak computational throughput of 85 GFLOP/sec. Mark the arithmetic intensity of the kernel on your plot and determine its expected performance.

## 2) [75 pts] Memory hierarchy: Matrix Vector Multiplication

A simple nested loop is sufficient to multiply an  $N \times M$  matrix with a  $M \times 1$  vector resulting in an  $N \times 1$  vector:

```
for (i = 0; i < N; i++) {
    for (j = 0; j < M; j++) {
        output[i] += input_matrix[i][j]*input_vector[j];
    }
}
```

However, as we discussed in class, this simple implementation can not make efficient use of the memory hierarchy. In particular, assuming row-major storage, while the accesses to the input matrix can achieve ideal cache miss rates, the accesses to the output matrix may result in high number of cache misses.

- a. [25 pts] *Implementation*: Following the cache blocking idea outlined in class, implement a cache optimized version of the matrix vector multiplication algorithm. For this purpose, you will use the source file provided which already includes the data structures you will use, an implementation of the simple multiplication algorithm, as well as time measurement mechanisms. Your task is to fill in the `optMultiplication` function.
- b. [20 pts] *Performance Analysis*: Test the performance of your implementation for a set of different  $N$  (number of rows),  $M$  (number of columns) and  $B$  (blocking factor) values. Identify the different regimes as indicated by the relative speed up of the optimized implementation over the naïve one, and outline the conditions (in terms of matrix sizes, shapes, etc.) under which you observe these regimes. **Hint**: Experiment with different matrix sizes (small to large) and input matrix shapes (short&wide rectangle, square, tall&skinny rectangle) while using a reasonable blocking factor.
- c. [20 pts] *Cache Performance Measurement*: Using a cache performance measurement tool (see instructions below on TAU) and your knowledge about how caches work, explain the root causes for the different regimes you observe in part b. For example, you can point to the ratio of cache misses you observe at various levels (L1, L2 or L3) using the naïve vs. optimized implementations to explain the relative performance differences. Can you make an estimation on the relative L1 and L2 latencies based on the performance and cache misses data you obtained?
- d. [15 pts] *Inference about the Memory Hierarchy*: Experiment with different blocking factors  $B$  to make inferences about the memory hierarchy. In particular, can you estimate the size of the L1 and L2 caches based on the performance and cache misses data you obtained? Compare your estimations with the hardware specifications.

### Instructions:

- **Using the dev-intel16 system:** You should use the dev-intel16 system for all of your runs. Whenever you login to your hpcc you are in gateway nodes where you should never run your programs. After login in hpcc run this command and then proceed with loading modules-

```
ssh dev-intel16
```

- **Compiling your programs:** While you can use any compiler (different versions of GCC or Intel compiler), throughout this assignment please use GCC/4.8.3. GCC/4.8.3 together with OpenMPI 1.8 are necessary for compiling and running with TAU, a performance analysis tool that we will be using. So whenever you login to HPCC or in your (batch or interactive) queue runs, make sure to execute the following commands:

```
module unload GNU OpenMPI
module load GNU/4.8.3 OpenMPI/1.8
```

After loading the proper modules, you can compile the provided source file in one of the two ways below.

```
make multiplication
make multiplication-tau
```

**Running the executables:** After compilation, the provided source can be executed either in the "test" or in the "perf" modes, conveniently on any of the dev-nodes:

```
./multiplication.x test inputfile
./multiplication.x perf N M B
```

where

inputfile contains a set of test runs. It starts with a line indicating the number of tests, followed by the tests themselves, each on a separate line. Each test is specified with N M B values.

N: number of matrix rows

M: number of matrix columns

B: blocking factor.

- We provide you a sample test file, which is the exact file that we will be using to test the accuracy (not performance) of your submission.
- **Measuring cache misses on HPCC:** For this purpose, we will be using the TAU performance analysis tool. Using hardware counters available through PAPI, TAU allows the measurement of an extensive list of CPU events. To compile your programs with TAU, follow the instructions below:

- i. Append the following lines to your Bash profile by editing the .bashrc file, which is a hidden file in your home directory (you can open it using vim or emacs, e.g. vim ~/.bashrc)

```
# PAPI
export PATH="/mnt/home/ohearnku/.local/papi/5.5.1/bin:$PATH"
export INCLUDE_PATH="/mnt/home/ohearnku/.local/papi/5.5.1/include:$INCLUDE_PATH"
export LD_LIBRARY_PATH="/mnt/home/ohearnku/.local/papi/5.5.1/lib:$LD_LIBRARY_PATH"
export MANPATH="/mnt/home/ohearnku/.local/papi/5.5.1/man:$MANPATH"

# TAU
export PATH="/mnt/home/ohearnku/.local/tau/2.26/x86_64/bin:$PATH"
export INCLUDE_PATH="/mnt/home/ohearnku/.local/tau/2.26/include:$INCLUDE_PATH"
```

```
export
LD_LIBRARY_PATH="/mnt/home/ohearnku/.local/tau/2.26/x86_64/lib:$LD_LIBRARY_PATH"
export TAU_MAKEFILE="/mnt/home/ohearnku/.local/tau/2.26/x86_64/lib/Makefile.tau-papi-
mpi"
export TAU_OPTIONS=-optCompInst
export TAU_METRICS="P_WALL_CLOCK_TIME:PAPI_L1_DCM:PAPI_L2_TCM:PAPI_L3_TCM"
```

- ii. To see the complete list of available hardware counters, run on a dev-node

```
papi_avail
```

where supported counters have "Yes" in the third and fourth columns.

- iii. Now, compile the source file using the Tau MPI wrapper:

```
make multiplication-tau
```

- iv. You can now run the executable *instrumented* with TAU on a dev-node or using the queue system:

```
mpiexec -np 1 tau_exec ./multiplication-tau 1000 1000000
```

200

- v. The performance data collected by TAU is saved automatically into default directories/files. Below are the directories/files you will see (as a result of the TAU\_METRICS options specified):

```
MULTI__P_WALL_CLOCK_TIME/profile.0.0.0
MULTI__PAPI_L1_DCM/profile.0.0.0
MULTI__PAPI_L2_TCM/profile.0.0.0
MULTI__PAPI_L3_TCM/profile.0.0.0
```

- vi. Normally, these files are intended to be viewed through a visualization software, for example using the paraprof command in the directory where the data was created:

```
paraprof &>/dev/null &
```

To view L1 data cache misses by function, click on the PAPI\_L1\_DCM item, and within the window it opens, click on either of the colored sections of the bar chart.

- vii. Paraprof requires X11, which can sometimes be tricky to run on a Windows or Mac, or if you do not have a very fast internet connection. We recommend simply viewing the profile.0.0.0 files for each event using the cat command or a text editor (vim or emacs). The contents are simple enough to be understood in this way.
- **Measuring your execution time and cache misses properly:** The wall-clock time measurement mechanism (based on the gettimeofday() function) implemented in the provided main.c file will allow you to measure the timings for a particular part of your program (see the skeleton code) precisely. However, on the dev-nodes there will be several other programs running simultaneously, and your measurements may not be very accurate due to the “noise”. After making sure that your program is bug-free and executes correctly (this can be done on your own computer if you have a C compiler

installed), a good way of getting reliable performance data for various input sizes is to use the interactive queue. **Please use the intel14 cluster for all your performance measurement runs!**

You can submit an interactive job request that will give you a dedicated node as follows:

```
qsub -I -l nodes=1:ppn=20,walltime=00:20:00,mem=64gb,feature=intel14
```

This will give you exclusive access to an intel14 node for 20 minutes. If you ask for a long job, your job may get delayed. **The default memory limit is set to be 750 MBs per job on HPC systems, so it is very important that you ask for more as above.**

Once you are granted an interactive job, make sure to run your jobs one after the other (i.e., do not run them as background jobs, and do not worry about it if you do not know what background jobs mean). This is important because having multiple background jobs running simultaneously may create “noise” in the data you obtain.

- **Batch jobs:** Interactive jobs may sometimes be delayed significantly, and therefore insisting on this option may be very counter-productive. For your convenience, we have provided a sample batch script named **job\_sample\_script.qsub**. Once you make sure that your program is running correctly on the dev-nodes, you can edit the batch script **run\_tests.bash** to submit multiple jobs together to the intel14 cluster. The job output will be located in the **tests/N** directory, where **N** is the test number. It is possible to set email alerts for when your job starts and ends (see the corresponding comments in the sample job script and replace with your e-mail address).
- **Obtaining files from the git repo & submission:** You will clone the skeleton code and the testing input file through the instructor repository on the Gitlab server. Assuming that you have already cloned the instructor repository, you will need to pull the most recently committed files for HW2 and move them over into your individual repository:

```
cd cse415-instructor
git pull
cd ../cse415-username/homework
cp -r ../../cse415-instructor/homework/2 .
```

Then complete the homework and submit it using your own personal repository. Your submission will include exactly two files:

- Your final `optMultiplication.c` file
- A pdf file named “HW2\_yourMSUNetID.pdf”, which contains your answers to the non-implementation questions of the assignment

*To submit your work, please follow the directions given in the “Homework Instructions” under the “Reference Material” section on D2L. Make sure to strictly follow these instructions; otherwise you may not receive proper credit.*