

CSE415: Introduction to Parallel Computing
Spring 2018, Homework 5

Deadline: April 20th, Friday 11:59 p.m

*This homework is 20% of your **homework grade**.*

Important note: Please use a word processing software (e.g., MS Word, Mac Pages, Latex, etc.) to type your homework. Follow the submission instructions at the end to turn in an electronic copy of your work.

1) [50 pts] Sqrt generation machine

Suppose you need to calculate square roots of the values in a huge array. Recently you have learnt distribute memory parallelization and these days you want to use distributed memory parallelization even in such a simple problem. The array will be created in the root node. You need to distribute the part of the array to different processes. But there is a small constraint. The distribution will not be of fixed length. Different processes will get different number of values of the array. For example, lets assume we have 4 processes and 100 numbers in our array. You are given this array –

21	26	23	30
----	----	----	----

This means that process 0 will get 21 numbers, process 1 will get 26 numbers and so on. Each process will receive their share of data from the root process and then they will calculate the square roots of each value they have and then they will send back their calculated values to the root node. In summary this should be the flow –

- Root node sends values according to the given array to the processes
- Each process computes square roots of their part.
- Each process then sends their calculated values back to the root node.

Note that, each process needs to allocate a space in their memory before receiving the data from the root process. For example, process 1 should receive a space for 26 elements before it receives the actual 26 elements from the process 0.

You can verify the results by printing first few numbers from the original array and the finally calculated array received in root 0.

You do not need to worry about the given array . The name of the given array is “proc_elems” in the code. Proc_elems will be generated for you. You have to use it for distribution.

Before compiling you code, make sure to load OpenMPI module. Then run the makefile provided.

Then run it using this command –

mpirun -n <process_count> ./mpi_sqrt.x <array size>

Try with 100mill, 500 mill, 1 bill array sizes with processes starting from 14 and its multiples upto 140. Each node in dev-intel16 has 28 cores. You need to allocate 5 nodes for your jobs.

2) [50 pts] Code Breaking

Cybersecurity is an important topic, and encryption (and how to crack it) is one of the most fundamental concepts in cybersecurity. Data Encryption Standard (DES) is a widely used encryption method that entails the use of a 56-bit secret key for encrypting and decrypting data. Better security standards such as the Advanced Encryption Standard (AES) that offer much wider key lengths has been gradually replacing DES since 2001.

In this problem, you will work on a brute-force cracking of encrypted messages using MPI. While DES-encryption can be broken in a matter of days using resources at HPCC, this is certainly not practical for our purposes. For this reason, we will look at a simple 32-bit encryption algorithm that relies on XOR operations along with bit rotations. The file `encrypter.c` implements this 32-bit encryption algorithm sequentially. Similarly, in file `codebreaker.c`, you are given a sequential code breaking algorithm which simply tries all 2^{32} keys in a brute force manner to decrypt a file encrypted by `encrypter.c`. While searching over the keys, the validation for each candidate is performed by comparing the words in the decrypted message against those in our small dictionary.

Your task is to parallelize the brute force algorithm given in `codebreaker.c` file in a few different ways, and analyze the performance of the resulting implementation. There are two general possibilities for partitioning this problem:

- **Static partitioning** where the division of the workload is done a priori based on a predetermined mechanism,
- **Dynamic partitioning** where the division and assignment of the workload is done at run-time, permitting load balancing to be performed, albeit at the cost of a more complicated solution.

We will only be looking at *static partitioning* techniques.

- Parallelize `codebreaker.c` using block partitioning of the set of possible keys. Your code should follow the following outline:
 - Root process should read the encrypted message and broadcast it to all processes,
 - Based on the total #processes, each process should determine its own set of trial keys,
 - When a process discovers the right key, it should print the success message and write the decrypted message into a file. In addition, it should notify all other processes so that the MPI program can be terminated immediately and correctly. **Hint:** This can be achieved by having each process set up an immediate receive operation prior to starting the loop for trying the keys. Then the process that discovers the key sends individual messages to all others to notify them about program completion. Clearly, all processes must be periodically (not very frequently though) testing on the immediate receive to see if any other process has had success.
 - Do not change the dictionary, the program output message, and/or the output file name.* Your program will be tested by providing a set of random (but properly encrypted) messages (of length < 1000) using the given dictionary. Two sample input files (`inp1.txt` and `inp2.txt`) are provided for your convenience.

Hint: You may find the commented out `printf` statements useful in understanding or debugging your parallel codes. **Make sure to comment out any `printf` statements that you uncomment before submission.** For initial testing, run your parallel codes with a small number of processes only, and try to decrypt messages encrypted with small key values (i.e. key = 10, 100, 1000).

- Analyze the scalability of your implementation in **part a** on a single node of the intel16 cluster. Note that as opposed to regular problems, search algorithms can achieve *super-linear*, *linear* or *zero* speedup compared to a sequential program (but the expected speed up will still be linear).

Demonstrate each of these possible scenarios by carefully choosing the encryption key and the number of processes (make sure to include examples using up to all 28 processes on a single intel16 cluster node).

- c) Identify and implement a partitioning strategy that overcomes the zero (or sublinear) speedup pitfall of the implementation in **part a**. Test version 2 of your codebreaker.c software on your examples from part b to show that your new partitioning strategy does indeed overcome the zero (or sublinear) speedup pitfall.

BONUS) [20 pts] MPI/OpenMP parallelization

On a multi-core architecture, good use of system resources can be achieved by first starting a single process for each physical socket in the system (on intel16, there are two sockets or chips per node), and then spawning as many threads as the number of cores per socket to fully exploit the multiple cores in that system. This is called MPI/OpenMP hybrid parallelization (obviously, when the threads created are OpenMP threads).

- Implement an MPI/OpenMP parallel version of the codebreaker.c program that you developed in question 3.
- Analyze the performance of the MPI/OpenMP parallel code with respect to the MPI only parallel version. Do you observe better or worse performance? Explain your observations.

Instructions:

- **Cloning git repo.** You can clone the skeleton codes through the git repo. Please refer to “*Homework Instructions*” under the “*Reference Material*” section on D2L.
- **Submission.** Your submission will include:
 - A pdf file named “HW5_yourMSUNetID.pdf”, which contains your answers to the non-implementation questions, and report and interpretation of performance results.
 - All source files used to generate the reported performance results. Make sure to use the exact files names listed below:
 - `Mpi_sqrt.c` (for the first problem)
 - `codebreaker.c` (for problem 2)
 - `codebreaker_v2.c` (for problem 2)
 - `codebreaker_hybrid.c` (for the bonus problem)

These are the default names in the git repository. It is essential that you do not change the directory structure or file names.

To submit your work, please follow the directions given in the “Homework Instructions” under the “Reference Material” section on D2L. Make sure to strictly follow these instructions; otherwise you may not receive proper credit.

- **Discussions.** For your questions about the homework, please use the Slack group for the class so that answers by the TA, myself (or one of your classmates) can be seen by others.
- **Compilation and execution.** You can the provided Makefile for compilation. Simply type “make” in the directory where the Makefile is located at.
- **Encryption and Decryption:**

You can encrypt a file using `encrypter.x`. running this command –

The resulting MPI parallel binary must be executed using `mpirun` (for example using 4 processes):

```
mpirun -n 4 ./encrypter.x filename_to_encrypt key
```

You can decrypt using this command –

```
mpirun -n 4 ./decrypter.x encrypted_filename
```

For hybrid MPI/OpenMP compilation in the bonus problem, you must provide the `-fopenmp` flag to the `mpicc` compiler (see the `CFLAGS` variable used by `mpicc` in the Makefile). Also remember to set the `OMP_NUM_THREADS` environment variable, when necessary.

- **Measuring your execution time properly.** The `MPI_Wtime()` or `omp_get_wtime()` commands will allow you to measure the timing for a particular part of your program (see the skeleton codes). *Make sure to collect 3-5 measurements and take their averages while reporting a performance data point.*
- **Executing your jobs.** On the dev-nodes there will be several other programs running simultaneously, and your measurements will not be accurate. After you make sure that your program is bug-free and executes correctly on the dev-nodes, the way to get good performance data for different programs and

various input sizes is to use the interactive or batch execution modes. *Note that jobs may wait in the queue to be executed for a few hours on a busy day, thus plan accordingly and do not wait for the last day of the assignment.*

- i) **Interactive queue.** Suggested options for starting an interactive queue on the **intel16** cluster is as follows:

```
qsub -I -l nodes=1:ppn=28,walltime=00:30:00,feature=intel16 -N myjob
```

The options above will allow exclusive access to a node for 30 minutes. If you ask for a long job, your job may get delayed. Note that default memory per job is 750 MBs, which should be plenty for the problems in this assignment. But if you will need more memory, you need to specify it in the job script.